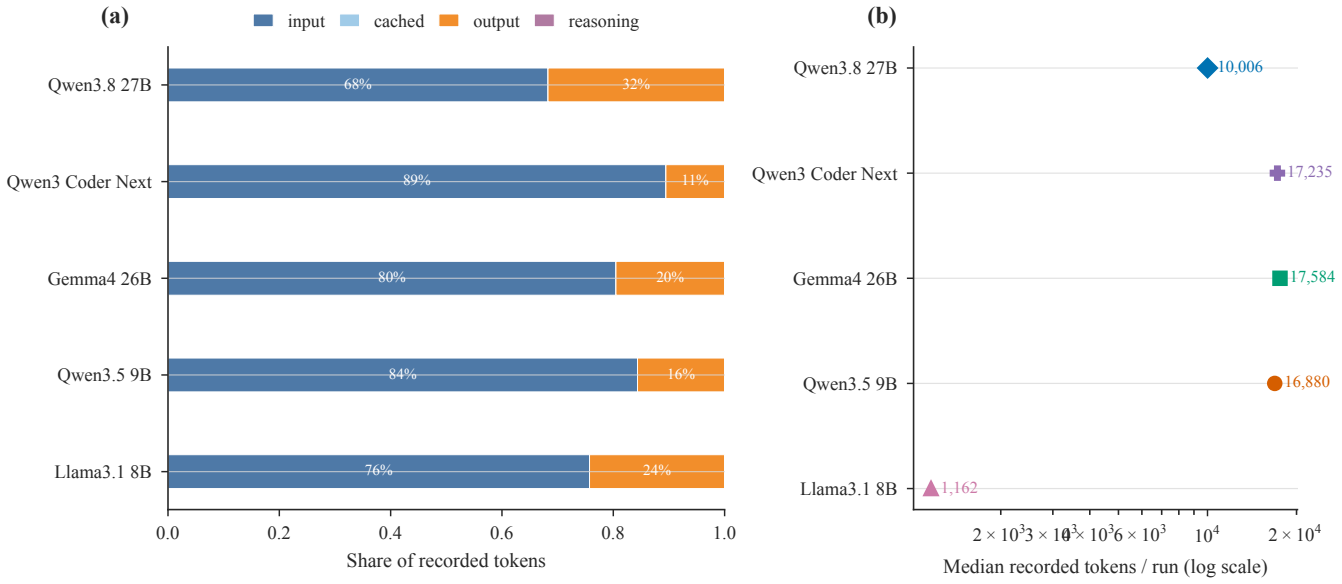
(a)
input
cached
output
reasoning
Qwen3.8 27B
Qwen3 Coder Next
Gemma4 26B
Qwen3.5 9B
Llama3.1 8B
68%
32%
89%
11%
80%
20%
84%
16%
76%
24%
0.0
0.2
0.4
0.6
0.8
1.0
Share of recorded tokens
(b)
Qwen3.8 27B
Qwen3 Coder Next
Gemma4 26B
Qwen3.5 9B
Llama3.1 8B
10,006
17,235
17,584
16,880
1,162
2×10^3
3×10^3
6×10^3
10^4
2×10^4
Median recorded tokens / run (log scale)