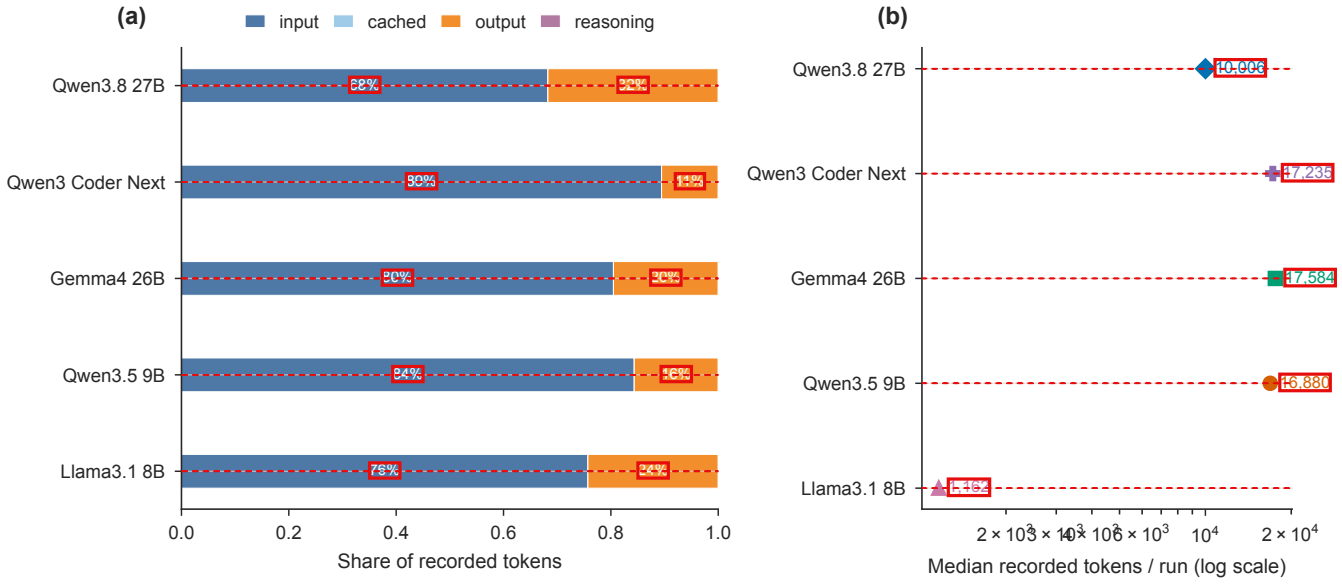
(a)
input
cached
output
reasoning
Qwen3.8 27B
Qwen3 Coder Next
Gemma4 26B
Qwen3.5 9B
Llama3.1 8B
68%
32%
80%
14%
80%
20%
84%
16%
76%
24%
0.0
0.2
0.4
0.6
0.8
1.0
Share of recorded tokens
(b)
Qwen3.8 27B
Qwen3 Coder Next
Gemma4 26B
Qwen3.5 9B
Llama3.1 8B
10,006
17,235
17,584
6,880
1,464
2 × 10³
3 × 10³
4 × 10³
6 × 10³
10⁴
2 × 10⁴
Median recorded tokens / run (log scale)